

SRI VISHNU K

Generative AI Engineer

Chennai, India | vishnukrishnakumar006@gmail.com | +91-7200409110 | [linkedin.com/in/srivishnuk](https://www.linkedin.com/in/srivishnuk)

PROFESSIONAL SUMMARY

Results-driven Generative AI Engineer with 4 years 9 months of IT experience, including 2.5+ years building production-grade LLM applications and RAG pipelines. Delivered enterprise AI solutions across manufacturing and financial services domains, processing 100,000+ documents and reducing manual review time by 60%. Proven expertise in end-to-end RAG pipeline development, AWS Bedrock (Claude), LangChain, LlamaIndex, and serverless AI architectures on AWS — with a solid foundation in Angular-based frontend development.

TECHNICAL SKILLS

Programming Languages: Python, SQL, JavaScript, TypeScript, C, C++, Bash

AI / Generative AI: LLMs, Prompt Engineering, RAG, Embeddings, Semantic Search, Text Summarization, Conversational AI, AI Workflow Automation

Frameworks & Libraries: LangChain, LlamaIndex, SentenceTransformers, Hugging Face Transformers, Angular, FastAPI, Pandas, NumPy, jQuery, AJAX

Vector Databases: Pinecone, FAISS

Cloud Platforms: AWS (Bedrock, Lambda, S3, API Gateway, Step Functions, EC2)

Frontend & CMS: HTML5, CSS3, Bootstrap, Responsive Design, WordPress, Elementor, SEO, Performance Tuning

APIs & Data Formats: RESTful APIs, JSON, XML

Tools & Platforms: Git, Docker (basic), GitHub Actions (CI/CD), Postman, VS Code, Jupyter Notebooks

WORK EXPERIENCE

Generative AI Engineer – Hotfoot Technology Solutions, Chennai, January 2024 – Present (2 Years 4 Months)

- Designed and deployed end-to-end RAG pipelines using **LangChain**, **LlamaIndex**, and **AWS Bedrock (Claude)**, enabling intelligent document Q&A and enterprise AI search across 2 major client engagements.
- Built and maintained RAG pipelines integrating **FAISS** and **Pinecone** vector databases with custom and cloud-based LLMs, improving query response accuracy by **40%** and reducing average search time from 8 minutes to under 30 seconds.
- Developed and maintained a versioned **prompt engineering** library (50+ prompt templates) using few-shot, chain-of-thought, and system-level instructions, reducing LLM hallucination rate by **35%**.
- Implemented structured logging and performance monitoring for AI pipelines — tracking token usage, latency, and retrieval quality — enabling iterative optimization that cut average response latency by **25%**.
- Architected serverless AI pipelines on **AWS** (Lambda, API Gateway, S3, Step Functions), handling 10,000+ daily API requests with 99.9% uptime across production environments.
- Built AI-powered document summarization tools that condensed 500-page policy documents, audit reports, and technical manuals into structured summaries, **reducing manual review time by 60%** across business units.
- Optimized LLM token usage through prompt compression, response caching, and intelligent chunking strategies, **reducing API costs by 30%** without compromising output quality.
- Integrated AI solutions with diverse enterprise data sources — PDFs, CSVs, SQL databases, and REST APIs — normalizing and preprocessing content for seamless ingestion into production AI pipelines.
- Collaborated cross-functionally with compliance, operations, and engineering teams across 2 enterprise clients in manufacturing and financial services domains.

Frontend Developer / Web Developer – Aptive Technologies, Chennai, August 2021 – January 2024 (2 Years 5 Months)

- Developed and maintained **15+ responsive web applications** using Angular, JavaScript, TypeScript, HTML5, and CSS3, delivering cross-browser compatible interfaces for diverse client projects.
- Built and customised **WordPress websites** using Elementor and Bootstrap, implementing SEO best practices that improved average page load speed by **40%** and increased organic search rankings for 3 clients.
- Integrated **RESTful APIs** and handled JSON/XML data exchange between frontend interfaces and backend services using AJAX and jQuery, supporting 5+ live production deployments.
- Collaborated with backend teams to define API contracts, ensuring seamless data flow and reducing integration bugs by **30%** through clear interface documentation.
- Applied Git-based version control and branching workflows across multi-developer projects, maintaining clean release histories for 10+ repositories.

PROJECTS

Enterprise Internal Knowledge Assistant (RAG on Custom LLM) | Jan 2024 – Aug 2024

Client: Confidential – Large Manufacturing Enterprise

Domain: Manufacturing / Operations Knowledge Management

Stack: Python, LangChain, FAISS, SentenceTransformers, REST APIs, AWS S3, AWS EC2

- Integrated a proprietary custom LLM with a secure RAG pipeline, enabling context-grounded responses from **50,000+ enterprise manufacturing documents** — reducing knowledge retrieval time by 70%.
- Designed custom chunking and embedding strategies using **SentenceTransformers** and FAISS, achieving a retrieval precision of 87% on domain-specific queries.
- Delivered measurable business impact: engineers reduced documentation search time from 20 minutes to under 3 minutes per query.

Enterprise AI Document Search & Compliance Assistant (AWS Bedrock + Claude) | Sep 2024 – Present

Client: Confidential – Financial Services / Banking Domain

Domain: Regulatory Compliance & Document Intelligence

Stack: AWS Bedrock (Claude), LlamaIndex, Pinecone, AWS Lambda, API Gateway, S3, Step Functions, Python

- Built a production-grade enterprise AI search and compliance assistant using **Claude via AWS Bedrock** and a **LlamaIndex + Pinecone** RAG pipeline, enabling compliance teams to query regulatory documents in natural language.
- Processed and indexed **10,000+ regulatory documents** across 5 compliance categories; reduced manual document review time by **60%** for a 20-person compliance team.
- Engineered serverless AWS architecture handling 500+ daily compliance queries with sub-2-second response times and full audit logging for regulatory traceability.
- Implemented output filtering, hallucination detection heuristics, and confidence scoring to ensure safe, compliant LLM responses in a regulated banking environment.

EDUCATION

Bachelor of Commerce (Information Technology)

2021 | CGPA: 8.75 / 10

Sri Krishna College of Arts and Science, Coimbatore

Higher Secondary Certificate (HSC)

2018 | 85%

Bharathiyar Higher Secondary School

Secondary School Leaving Certificate (SSLC)

2016 | 95%

Bharathiyar Matriculation School